

# DS4001-25SP-HW2：搜索

刘芮希 PB22010402

2025 年 5 月 8 日

## 1 问题 1：马尔可夫决策过程[9%=6%+3%]

(a) 迭代计算

值迭代核心公式为

$$V_{k+1}(s) = \max_a \left[ \sum_{s'} P(s' | s, a) (R(s, a, s') + \gamma V_k(s')) \right]$$

下面进行数值计算

**初始值：**  $V^{(0)}(-2) = 0, V^{(0)}(-1) = 0, V^{(0)}(0) = 0, V^{(0)}(1) = 0, V^{(0)}(2) = 0$

i=0 到 i=1:

状态 -1: 动作  $a_1$  的 Q 值为  $0.7 \times 10 + 0.2 \times (-1) + 0.1 \times (-1) = 6.7$ , 动作  $a_2$  的 Q 值为  $0.5 \times 10 + 0.3 \times (-1) + 0.2 \times (-1) = 4.5$ , 取最大值 6.7。

状态 0: 无论动作  $a_1$  或  $a_2$ , Q 值均为 -1.0 (转移概率加权后均为负值), 取 -1.0。

状态 1: 动作  $a_1$  的 Q 值为  $0.2 \times 20 + 0.7 \times (-1) + 0.1 \times (-1) = 3.2$ , 动作  $a_2$  的 Q 值为  $0.3 \times 20 + 0.5 \times (-1) + 0.2 \times (-1) = 5.3$ , 取最大值 5.3。

则  $V^{(1)}(-2) = 0, V^{(1)}(-1) = 6.7, V^{(1)}(0) = -1.0, V^{(1)}(1) = 5.3, V^{(1)}(2) = 0$

i=1 到 i=2:

状态 -1: 动作  $a_1$  的 Q 值为  $0.710 + 0.2(-1 + V^{(1)}(0)) + 0.1(-1 + V^{(1)}(-1)) = 7.17$ , 动作  $a_2$  的 Q 值为  $0.510 + 0.3(-1 + V^{(1)}(0)) + 0.2(-1 + V^{(1)}(-1)) = 5.54$ , 取最大值 7.17。

状态 0: 动作  $a_1$  的 Q 值为  $0.2(-1 + V^{(1)}(1)) + 0.1(-1 + V^{(1)}(0)) + 0.7(-1 + V^{(1)}(-1)) = 4.65$ , 动作  $a_2$  的 Q 值为  $0.3(-1 + V^{(1)}(1)) + 0.2(-1 + V^{(1)}(0)) + 0.5(-1 + V^{(1)}(-1)) = 3.74$ , 取最大值 4.65。

状态 1: 动作  $a_1$  的 Q 值为  $0.220 + 0.1(-1 + V^{(1)}(1)) + 0.7(-1 + V^{(1)}(0)) = 3.03$ , 动作  $a_2$  的 Q 值为  $0.320 + 0.2(-1 + V^{(1)}(1)) + 0.5(-1 + V^{(1)}(0)) = 5.86$ , 取最大值 5.86。

则  $V^{(2)}(-2) = 0, V^{(2)}(-1) = 7.17, V^{(2)}(0) = 4.65, V^{(2)}(1) = 5.86, V^{(2)}(2) = 0$

(b)  $i = 2$  时最优策略

$s = -1$ 时, 动作  $a_1$  的 Q 值为7.17, 动作  $a_2$  的 Q 值为5.54, 最优策略 $\mu(-1) = a_1$

$s = 0$ 时, 动作  $a_1$  的 Q 值为4.65, 动作  $a_2$  的 Q 值为3.74, 最优策略 $\mu(0) = a_1$

$s = 1$ 时, 动作  $a_1$  的 Q 值为3.03, 动作  $a_2$  的 Q 值为5.86, 最优策略 $\mu(1) = a_2$

$(s, \mu(s))$  数值对为  $(-1, a_1), (0, a_1), (1, a_2)$

## 2 问题 2: Q-Learning[12%=3%+6%+3%]

(a) 推导  $Q(s, a)$

动作价值函数的定义为  $Q(s, a) = \mathbb{E}[G_t \mid s_t = s, a_t = a]$ , 其中  $G_t = \sum_{k=t}^{\infty} \gamma^{k-t} R_k$

执行动作  $a$  后, 环境转移到状态  $s'$ , 获得即时奖励  $R(s, a, s')$ , 后续回报为  $\gamma G_{t+1}$ 。因此:

$$\mathbb{E}[G_t \mid s_t = s, a_t = a] = \mathbb{E}[R(s, a, s') + \gamma G_{t+1} \mid s_t = s, a_t = a]$$

对所有可能的下一状态  $s'$ , 计算其期望:

$$\mathbb{E}[R(s, a, s') + \gamma G_{t+1} \mid s_t = s, a_t = a] = \sum_{s'} P(s' \mid s, a) [R(s, a, s') + \gamma \mathbb{E}[G_{t+1} \mid s_{t+1} = s']]$$

根据状态价值函数的定义  $V(s') = \mathbb{E}[G_{t+1} \mid s_{t+1} = s']$  代入得:

$$\sum_{s'} P(s' \mid s, a) [R(s, a, s') + \gamma \mathbb{E}[G_{t+1} \mid s_{t+1} = s']] = \sum_{s'} P(s' \mid s, a) (R(s, a, s') + \gamma V(s'))$$

故  $Q(s, a) = \sum_{s'} P(s' \mid s, a) (R(s, a, s') + \gamma V(s'))$

(b) Monte-Carlo 更新过程

初始值  $q(s, a) = 0$ , 折扣因子  $\gamma = 1$ , 学习率  $\alpha = 1$ 。轨迹为:

$$\tau = \{(0, a_1, 2), (1, a_1, 3), (0, a_2, -1), (1, a_2, 0)\}$$

$t = 1$

轨迹片段:  $(s = 0, a = a_1, r = 2)$

计算回报  $G_1$ :

$$G_1 = 2 + 3 + (-1) + 0 = 4$$

更新动作价值函数:

$$q(0, a_1) \leftarrow G_1 = 4$$

更新后的  $q$ -表:

$s$	$a_1$	$a_2$
0	4	0
1	0	0

$t = 2$

轨迹片段:  $(s = 1, a = a_1, r = 3)$

计算回报  $G_2$ :

$$G_2 = 3 + (-1) + 0 = 2$$

更新动作价值函数:

$$q(1, a_1) \leftarrow G_2 = 2$$

更新后的  $q$ -表:

$s$	$a_1$	$a_2$
0	4	0
1	2	0

$t = 3$

轨迹片段:  $(s = 0, a = a_2, r = -1)$

计算回报  $G_3$ :

$$G_3 = -1 + 0 = -1$$

更新动作价值函数:

$$q(0, a_2) \leftarrow G_3 = -1$$

更新后的  $q$ -表:

$s$	$a_1$	$a_2$
0	4	-1
1	2	0

$t = 4$

轨迹片段:  $(s = 1, a = a_2, r = 0)$

计算回报  $G_4$ :

$$G_4 = 0$$

更新动作价值函数:

$$q(1, a_2) \leftarrow G_4 = 0$$

最终  $q$ -表:

$s$	$a_1$	$a_2$
0	4	-1
1	2	0

各时间步更新后的动作价值函数值:

$$t = 1: \quad q(0, a_1) = 4$$

$$t = 2: \quad q(1, a_1) = 2$$

$$t = 3: \quad q(0, a_2) = -1$$

$$t = 4: \quad q(1, a_2) = 0$$

### (c) Q-Learning 收敛性核心思路

Q-Learning 算法能够收敛的核心原因在于其通过 贝尔曼最优算子的收缩性 和 状态-动作对的充分探索, 逐步缩小估计值与最优值之间的误差, 最终收敛到最优动作价值函数  $Q^*$ 。

#### 1. 贝尔曼最优算子的收缩性

Q-Learning 的更新本质是反复应用 贝尔曼最优算子  $\mathcal{T}$ , 其定义为:

$$(\mathcal{T}Q)(s, a) = R(s, a) + \gamma \max_{a'} Q(s', a')$$

该算子具有 **压缩映射** 性质，即对任意两个函数  $Q_1$  和  $Q_2$ ，满足：

$$\|\mathcal{T}Q_1 - \mathcal{T}Q_2\|_\infty \leq \gamma \|Q_1 - Q_2\|_\infty$$

由于  $\gamma \in [0, 1)$ ，每次应用算子后，函数之间的最大差异会以  $\gamma$  的倍数衰减。反复迭代后， $Q$  值会逐渐逼近唯一的不动点  $Q^*$ ，即最优动作价值函数。

## 2. 误差的指数衰减

假设初始  $Q$  表与  $Q^*$  的最大误差为  $\Delta_0$ ，每次更新后，误差会被压缩为  $\gamma\Delta_0$ 。在 **确定性环境** 中，若每个状态-动作对  $(s, a)$  被 **无限次访问**，经过  $k$  轮遍历所有状态-动作对后，最大误差降至  $\gamma^k\Delta_0$ 。当  $k \rightarrow \infty$ ，误差趋近于零， $Q$  表收敛到  $Q^*$ 。

## 3. 步长条件与推广

对于步长  $\alpha \in (0, 1]$ ，更新公式可视为当前估计值与贝尔曼目标的加权平均：

$$Q_{n+1}(s, a) = (1 - \alpha)Q_n(s, a) + \alpha\mathcal{T}Q_n(s, a)$$

在满足 **无限次访问条件** 时，误差上界变为  $(\alpha\gamma + 1 - \alpha)^k\Delta_0$ 。由于  $\alpha\gamma + (1 - \alpha) < 1$ ，误差仍以指数速率衰减，保证收敛。

## 4. 收敛前提条件

- **有限 MDP**：状态与动作空间有限。
- **充分探索**：每个状态-动作对被无限次访问（如通过  $\epsilon$ -贪心策略）。
- **确定性环境**：状态转移和奖励函数确定（可推广至非确定性环境，但需额外条件）。

Q-Learning 的收敛性依赖于 **贝尔曼算子的压缩性** 和 **状态-动作对的充分探索**，通过迭代逐步消除估计误差，最终逼近最优策略。其本质是通过动态规划的收缩过程，结合增量式更新，实现对长期回报的准确估计。

# 3 问题 3: Gobang Programming[53%=33%+10%+10%]

(a) [代码]

```
1 def get_next_state(self, action: Tuple[int, int, int], noise: Tuple[int, int, int]) -> np.array:
2     next_state = self.board.copy()
3     piece, x, y = action
4     next_state[x][y] = piece
5
6     if noise is not None:
7         white, x_white, y_white = noise
8         next_state[x_white][y_white] = white
9     return next_state

```

```
1 def sample_noise(self) -> Union[Tuple[int, int, int], None]:
2     if self.action_space:
3         x, y = random.choice(self.action_space)
4         self.action_space.remove((x, y))
5         return 2, x, y
6     else:
7         return None

```

```
1 def get_connection_and_reward(self, action: Tuple[int, int, int],
2 noise: Tuple[int, int, int]) -> Tuple[int, int, int, int, float]:
3     black_1, white_1 = self.count_max_connections(self.board)
4     next_state = self.get_next_state(action, noise)
5     black_2, white_2 = self.count_max_connections(next_state)
6     reward = (black_2 - black_1) - (white_2 - white_1)
7     return black_1, white_1, black_2, white_2, reward

1 def sample_action_and_noise(self, eps: float):
2     s = self.array_to_hashable(self.board)
3     if random.random() < eps or s not in self.Q:
4         x, y = random.choice(self.action_space)
5         action = (1, x, y)
6     else:
7         valid_actions = [(a, q) for a, q in self.Q[s].items() if (a[1], a[2]) in self.action_space]
8         if valid_actions:
9             action = max(valid_actions, key=lambda item: item[1])[0]
10        else:
11            x, y = random.choice(self.action_space)
12            action = (1, x, y)
13        self.action_space.remove((action[1], action[2]))
14        return action, self.sample_noise()

1 def q_learning_update(self, s0_: np.array, action: Tuple[int, int, int],
2 s1_: np.array, reward: float, alpha_0: float = 1):
3     s0, s1 = self.array_to_hashable(s0_), self.array_to_hashable(s1_)
4     self.s_a_visited[(s0, action)] = 1 if (s0, action) not in self.s_a_visited else \
5         self.s_a_visited[(s0, action)] + 1
6     alpha = alpha_0 / self.s_a_visited[(s0, action)]
7     if s0 not in self.Q:
8         self.Q[s0] = {}
9     current_q = self.Q[s0].get(action, 0.0)
10    if s1 in self.Q and self.Q[s1]:
11        max_q_next = max(self.Q[s1].values())
12    else:
13        max_q_next = 0.0
14    new_q = current_q + alpha * (reward + max_q_next - current_q)
15    self.Q[s0][action] = new_q
```

(b)  $n=3$  训练及测试结果

```
(ds) PS E:\USTC-DS4001-25sp\Homework\HW2\code> & D:\Anaconda\envs\ds\python.exe e:\USTC-DS4001-25sp\Homework\HW2\code\learner.py
100%| 10000/10000 [00:30<00:00, 325.94it/s]
learning ended.
```

图 1:  $n=3$  训练

```
The evaluated winning probability for the black pieces is 0.9655870445344129.
Black wins: 955, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9656218402426694.
Black wins: 956, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9656565656565657.
Black wins: 957, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9656912209889001.
Black wins: 958, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9657258064516129.
Black wins: 959, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9657603222557906.
Black wins: 960, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.965794768661167.
Black wins: 961, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9658291457286432.
Black wins: 962, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9658634538152611.
Black wins: 963, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9658976930792377.
Black wins: 964, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.9659318637274549.
Black wins: 965, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.965965965965966.
Black wins: 966, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.966.
100%
Evaluation finished. Black wins: 966, white wins: 5, and ties: 29.
The evaluated winning probability for the black pieces is 0.966.
1000/1000 [00:00:00:00, 1886.08it/s]
```

图 2:  $n=3$ 测试

(c)  $n=4$  训练及测试结果

```
(ds) PS E:\USTC-DS4001-25sp\Homework\HW2\code> & D:\Anaconda\envs\ds\python.exe e:/USTC-DS4001-25sp/Homework/HW2/code/learner.py  
100%|██████████████████████████████████████████████████████████████████████████| 10000/10000 [01:30<00:00, 110.45it/s]  
learning ended.
```

图 3:  $n=4$ 训练

```
The evaluated winning probability for the black pieces is 0.6690283408089717.
Black wins: 662, white wins: 114, and ties: 213.
The evaluated winning probability for the black pieces is 0.6693629929221436.
Black wins: 663, white wins: 114, and ties: 213.
The evaluated winning probability for the black pieces is 0.6696969696969697.
Black wins: 664, white wins: 114, and ties: 213.
The evaluated winning probability for the black pieces is 0.6700302724520686.
Black wins: 664, white wins: 115, and ties: 213.
The evaluated winning probability for the black pieces is 0.669354837096774.
Black wins: 665, white wins: 115, and ties: 213.
The evaluated winning probability for the black pieces is 0.6696878147029205.
Black wins: 665, white wins: 115, and ties: 214.
The evaluated winning probability for the black pieces is 0.6690140845078423.
Black wins: 666, white wins: 115, and ties: 214.
The evaluated winning probability for the black pieces is 0.6693467336683417.
Black wins: 666, white wins: 115, and ties: 215.
The evaluated winning probability for the black pieces is 0.6686746987951807.
Black wins: 667, white wins: 115, and ties: 215.
The evaluated winning probability for the black pieces is 0.6690070210631895.
Black wins: 668, white wins: 115, and ties: 215.
The evaluated winning probability for the black pieces is 0.6693386773547094.
Black wins: 669, white wins: 115, and ties: 215.
The evaluated winning probability for the black pieces is 0.6696969696969697.
Black wins: 670, white wins: 115, and ties: 215.
The evaluated winning probability for the black pieces is 0.67.
100%
Evaluation finished. Black wins: 670, white wins: 115, and ties: 215.
The evaluated winning probability for the black pieces is 0.67.
1000/1000 [00:01:00:00, 789.751t/s]
```

图 4:  $n=4$ 测试

#### 4 问题 4: Deeper Understanding[16%=5%+5%+2%+4%]

### 4.1 Bellman 算子与压缩映射

(a) 证明 Bellman 算子是压缩映射

### 1. 定义 Bellman 算子:

$$(\mathcal{T}v)(s) = \max_{a \in A} \left\{ r_{sa} + \gamma \cdot \sum_{s' \in S} p_{sas'} \cdot v(s') \right\}.$$

对于两个价值函数  $v_1$  和  $v_2$ ，分别有：

$$(\mathcal{T}v_1)(s) = \max_{a \in A} \left\{ r_{sa} + \gamma \cdot \sum_{s' \in S} p_{sas'} \cdot v_1(s') \right\},$$

$$(\mathcal{T}v_2)(s) = \max_{a \in A} \left\{ r_{sa} + \gamma \cdot \sum_{s' \in S} p_{sas'} \cdot v_2(s') \right\}.$$

2. 构造差值：对于任意状态  $s$ ，设  $a_1^*$  和  $a_2^*$  分别是使得  $\mathcal{T}v_1$  和  $\mathcal{T}v_2$  达到最大值的动作：

$$(\mathcal{T}v_1)(s) = r_{sa_1^*} + \gamma \cdot \sum_{s' \in S} p_{sa_1^*s'} \cdot v_1(s'),$$

$$(\mathcal{T}v_2)(s) = r_{sa_2^*} + \gamma \cdot \sum_{s' \in S} p_{sa_2^*s'} \cdot v_2(s').$$

由于  $a_1^*$  是  $v_1$  下的最优动作，而  $a_2^*$  是  $v_2$  下的最优动作，因此：

$$(\mathcal{T}v_1)(s) \geq r_{sa_2^*} + \gamma \cdot \sum_{s' \in S} p_{sa_2^*s'} \cdot v_1(s'),$$

$$(\mathcal{T}v_2)(s) \geq r_{sa_1^*} + \gamma \cdot \sum_{s' \in S} p_{sa_1^*s'} \cdot v_2(s').$$

利用这些不等式，可以得到：

$$(\mathcal{T}v_1)(s) - (\mathcal{T}v_2)(s) \leq \gamma \cdot \sum_{s' \in S} p_{sa_2^*s'} (v_1(s') - v_2(s')) \leq \gamma \|v_1 - v_2\|_\infty,$$

$$(\mathcal{T}v_2)(s) - (\mathcal{T}v_1)(s) \leq \gamma \cdot \sum_{s' \in S} p_{sa_1^*s'} (v_2(s') - v_1(s')) \leq \gamma \|v_1 - v_2\|_\infty.$$

因此：

$$|(\mathcal{T}v_1)(s) - (\mathcal{T}v_2)(s)| \leq \gamma \|v_1 - v_2\|_\infty.$$

3. 取最大值：对所有状态  $s$  取最大值，得到：

$$\|\mathcal{T}v_1 - \mathcal{T}v_2\|_\infty = \max_{s \in S} |(\mathcal{T}v_1)(s) - (\mathcal{T}v_2)(s)| \leq \gamma \|v_1 - v_2\|_\infty.$$

综上，Bellman 算子  $\mathcal{T}$  是最大范数下的  $\gamma$ -压缩映射。

(b) 证明最多一个不动点

假设存在两个不同的不动点  $v_1$  和  $v_2$ ，即：

$$\mathcal{T}v_1 = v_1, \quad \mathcal{T}v_2 = v_2.$$

根据  $\gamma$ -压缩性质，有：

$$\|v_1 - v_2\|_\infty = \|\mathcal{T}v_1 - \mathcal{T}v_2\|_\infty \leq \gamma \|v_1 - v_2\|_\infty.$$

整理得：

$$(1 - \gamma) \|v_1 - v_2\|_\infty \leq 0.$$

由于  $\gamma \in [0, 1)$ , 故  $1 - \gamma > 0$ , 这意味着:

$$\|v_1 - v_2\|_\infty \leq 0.$$

而范数具有非负性, 因此:

$$\|v_1 - v_2\|_\infty = 0 \quad \Rightarrow \quad v_1 = v_2.$$

这与假设  $v_1$  和  $v_2$  不同矛盾, 故  $\mathcal{T}$  最多只能有一个不动点。

## 4.2 重要性采样

### (a) 重要性采样等式证明

假设分布  $q(x)$  的支撑集包含  $p(x)$  的支撑集, 即当  $p(x) > 0$  时,  $q(x) > 0$ 。若此条件不满足, 则  $\frac{p(x)}{q(x)}$  在  $q(x) = 0$  时无定义, 需额外处理。

对于  $x \sim p$  的期望:

$$\mathbb{E}_{x \sim p}[f(x)] = \int f(x) p(x) dx \quad (\text{连续情况}) \quad (1)$$

$$\text{或 } \mathbb{E}_{x \sim p}[f(x)] = \sum_x f(x) p(x) \quad (\text{离散情况}). \quad (2)$$

对右侧的  $x \sim q$  的期望进行展开:

$$\mathbb{E}_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right] = \int f(x) \frac{p(x)}{q(x)} q(x) dx \quad (\text{连续情况}) \quad (3)$$

$$\text{或 } \mathbb{E}_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right] = \sum_x f(x) \frac{p(x)}{q(x)} q(x) \quad (\text{离散情况}). \quad (4)$$

在两种情况下,  $q(x)$  与分母的  $q(x)$  抵消:

$$\int f(x) p(x) dx = \mathbb{E}_{x \sim p}[f(x)] \quad (\text{连续情况}), \quad (5)$$

$$\sum_x f(x) p(x) = \mathbb{E}_{x \sim p}[f(x)] \quad (\text{离散情况}). \quad (6)$$

在支撑集条件满足时, 有:

$$\mathbb{E}_{x \sim p}[f(x)] = \mathbb{E}_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right].$$

### (b) 方差公式证明

根据方差公式  $\text{Var}[X] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$ :

$$\text{Var}_{x \sim p}[f(x)] = \mathbb{E}_{x \sim p}[f(x)^2] - (\mathbb{E}_{x \sim p}[f(x)])^2,$$

$$\text{Var}_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right] = \mathbb{E}_{x \sim q} \left[ \left( f(x) \frac{p(x)}{q(x)} \right)^2 \right] - \left( \mathbb{E}_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right] \right)^2.$$

由重要性采样性质  $\mathbb{E}_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right] = \mathbb{E}_{x \sim p}[f(x)]$ , 可得:

$$\left( \mathbb{E}_{x \sim q} \left[ f(x) \frac{p(x)}{q(x)} \right] \right)^2 = (\mathbb{E}_{x \sim p}[f(x)])^2.$$



代入方差表达式

$$\begin{aligned}\mathrm{Var}_p[f(x)] - \mathrm{Var}_q\left[f(x)\frac{p(x)}{q(x)}\right] &= \left[\mathbb{E}_p[f(x)^2] - (\mathbb{E}_p[f(x)])^2\right] \\ &\quad - \left[\mathbb{E}_q\left[f(x)^2\frac{p(x)^2}{q(x)^2}\right] - (\mathbb{E}_p[f(x)])^2\right] \\ &= \mathbb{E}_p[f(x)^2] - \mathbb{E}_q\left[f(x)^2\frac{p(x)^2}{q(x)^2}\right].\end{aligned}$$

利用重要性采样：

$$\mathbb{E}_q\left[f(x)^2\frac{p(x)^2}{q(x)^2}\right] = \mathbb{E}_p\left[f(x)^2\frac{p(x)}{q(x)}\right].$$

故

$$\mathrm{Var}_p[f(x)] - \mathrm{Var}_q\left[f(x)\frac{p(x)}{q(x)}\right] = \mathbb{E}_p[f(x)^2] - \mathbb{E}_p\left[f(x)^2\frac{p(x)}{q(x)}\right].$$

## 体验反馈[10%]

- (a) [必做] 8h
- (b) [选做] 我五子棋能下得过AI吗